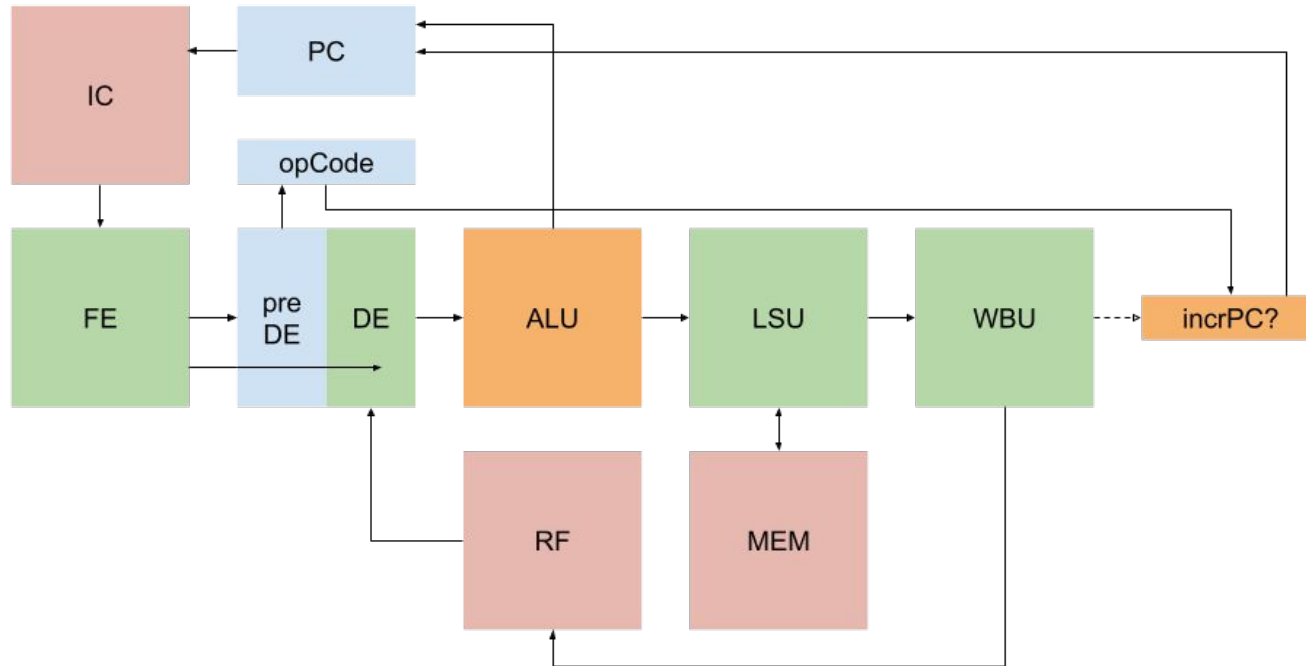



# Processor Architecture



# Instruction Set Architecture

add	rd	rs	rt	Reg[rd] = Reg[rs] + Reg[rt]
addi	rd	rs	#n	Reg[rd] = Reg[rs] + n
mul	rd	rs	rt	Reg[rd] = Reg[rs] + Reg[rt]
muli	rd	rs	#n	Reg[rd] = Reg[rs] + n
cmp	rd	rs	rt	Reg[rd] = rs < rs ? -1 : rs == rs ? 0 : 1
ld	rd	ro	#n	Reg[rd] = Mem[Reg[ro] + n]
ldc	rd	#n		Reg[rd] = Mem[n]
st	rs	ro	#n	Mem[Reg[ro] + n] = rf
brlz	rs	#n		rf <= 0 ? Pc = n : no-op
jplz	rs	#n		rf <= 0 ? Pc = Pc + n : no-op
br	#n			Pc = n
jp	#n			Pc = Pc + n

destination register ::= rd

source register ::= rs | rt

immediate ::= #n

## registers

*zero*, //constant zero

*ra*, //return address

*sp*, //stack pointer

*gp*, //global pointer

*tp*, //thread pointer

*t0*, *t1*, *t2*, *t3*, *t4*, *t5*, *t6*, // temporary registers

*s0*, *s1*, *s2*, *s3*, *s4*, *s5*, *s6*, *s7*, *s8*, *s9*, *s10*, *s11*, //saved

*a0*, *a1*, *a2*, *a3*, *a4*, *a5*, *a6*, *a7* //argument registers

## example code

```
loop:      add s6 s5 t0      --loc + i
           ld s6 s6 #0      --mem[loc + i]
           add s7 s5 t0
           addi s7 s7 #1    --loc + i + 1
           ld s7 s7 #0      --mem[loc + i + 1]
           cmp s8 s7 s6     -- s6 < s7 ?
           addi s8 s8 #1     --
           add s0 s5 t0     -- loc + i
           addi s1 s0 #1     -- loc + i + 1
           brlz s8 swap     -- labels
```

# Planned experiments

- Non-pipelined versus pipelined speedup
  - Hypothesis: without dependencies IPC will be proportional to pipeline depth/number of execution units
- In-order versus OoO speedup
  - Hypothesis: notable reduction in number of cycles from better utilisation of execution units
- Branch predictor accuracy for performance for programs with repetition
  - Hypothesis: more sophisticated/higher number of states for predictor will improve accuracy of predictions. Furthermore this gives notable reduction in cycles for the same programs
- Stall reduction with register renaming
  - Hypothesis: notable reduction in false dependencies with register renaming compared to without
- Average execution unit utilisation
  - Hypothesis: higher superscalar issue width proportionally increases execution utilisation until plateau when execution units are saturated